

Orthogonal Subspaces

4 Matrix Spaces:

1) $N(\underline{A})$ is the orthogonal complement of $C(\underline{A}^T)$ in \mathbb{R}^m ,

2) $N(\underline{A}^T)$ is the orthogonal complement of $C(\underline{A})$

Orthogonal complement means every vector in one space is orthogonal to all vectors in the other space,

$$\underline{A}\underline{x} = \underline{b} \quad \text{let } \underline{y} \in N(\underline{A})$$

$$\Rightarrow \underline{A}\underline{x} + \underline{A}\underline{y} = \underline{b}$$

Approximation & Round-off Error

- In many cases, no analytic solution exists
- Numerical solutions provide approximate results that should be close to the true result.
- But: we can not typically compute the errors of the numerical method,
 - Rarely is the input exact
 - Algorithms introduce errors
 - The results thus depend on both of these errors.


The question is this:

"How much error is present in our calculation & is that error tolerable?"

Highly Critical: Identification, Quantification & minimization of error,

Accuracy vs Precision

- Accuracy: How close the value is to the true value, typically thought of "How accurate is the method?"
- Precision: How closely do the computed values agree.



Numerical errors:
Errors arise from numerical approximations,
model choice, system, etc.

$$\text{Truth} = \text{approximate Value} + \text{Error}$$
$$x^* = x + e$$

$\Rightarrow e = x^* - x$ The error,

Relative Error: $\frac{\text{Truth} - \text{approx}}{\text{truth}}$

$$e_{rel} = \frac{x^* - x}{x^*}, \quad e_{rel\%} = e_{rel} * 100\%$$

Typically we can't compute the error. Why?

You need the true solution!


If you have that, why approximate?

Approximate error via:

$$e_{approx} = \frac{x^{new} - x^{old}}{x^{old}}$$

Significant Figures

Look at this gauge!



$x = 26.5?$
 $= 26.55?$

Significant Figures tell how much useful information do we really have,

Typically the digits that we are certain of + one more,

ex.) 53,200

5.32×10^4	3 sig digits
5.320×10^4	4
5.3200×10^4	5

ex.) leading zeros do not count towards this.

0.001753	>	All have 4 sig. figures
0.01753		
0.1753		

Sig figs tell us which numbers you can use with confidence,

- 32-bit machine : 7 significant figures.
- 64-bit machine : 15 significant figures.

- Double Precision: more precise than single precision, but slower to use.

$$\pi = 3,14159265358979 \mid 3238462643 \dots$$

←

False Significant Figures

ex1) $\frac{3,25 \times}{1,96 \times} = 1,65816326 \dots$ (from matlab)

In practice, only report 1,65 (Chopping) or 1,66 (Rounding)

Why?

Because in reality we do not know the third decimal point value,

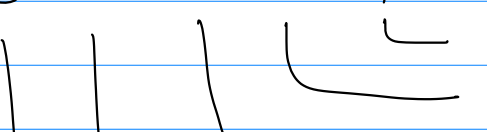
ex1) $\frac{3,259}{1,960} = 1,662755 \dots$ > Chopping
 $\frac{3,250}{1,969} = 1,65058 \dots$

$\frac{3,254}{1,955} = 1,664415 \dots$ > Rounding
 $\frac{3,245}{1,964} = 1,6522 \dots$

Numbering System

Base 10 (Decimals) $0, \dots, 9$

8 6 4 0 9



9	$\times 10^0$	=	9
0	$\times 10^1$	=	0
4	$\times 10^2$	=	400
6	$\times 10^3$	=	6000
8	$\times 10^4$	=	80000
<hr/>			
86409			

Base-2 (Binary) 0, 1

1 0 1 0 1 1 0 1

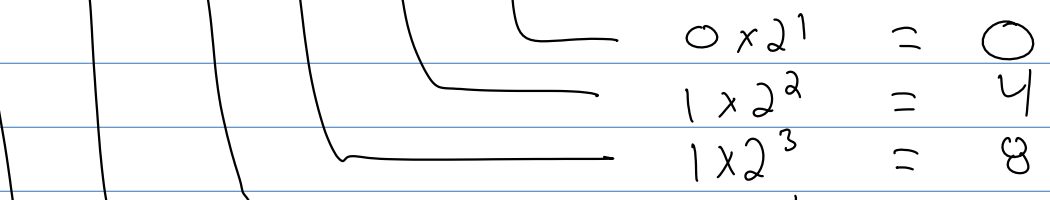


Diagram illustrating the binary representation of the decimal number 173 using powers of 2. The diagram shows a sequence of L-shaped curves, each representing a power of 2, and a corresponding list of powers of 2 with their decimal values.

Power of 2	Decimal Value
1×2^0	1
0×2^1	0
1×2^2	4
1×2^3	8
0×2^4	0
1×2^5	32
0×2^6	0
1×2^7	128

The sum of the decimal values is 173 (in base-10).

Other Systems

Base-8 (Octal) 0, 1, 2, 3, 4, 5, 6, 7

Base-16 (Hexadecimal) : 0, 1, ..., 15

$$\text{base-10: } 0.3125 = 3 \times 10^{-1} + 1 \times 10^{-2} + 2 \times 10^{-3} + 5 \times 10^{-4}$$

$$\text{base 2: } 101101 : 1 \times 2^5 + 0 \times 2^4 + 1 \times 2^3 + 1 \times 2^2 + 0 \times 2^1 + 1 \times 2^0 \\ = 45$$

$$0.1011 = 1 \times 2^{-1} + 0 \times 2^{-2} + 1 \times 2^{-3} + 1 \times 2^{-4} = \frac{11}{16}$$


Finite Precision Arithmetic

- One source of error is the limitation of the computer to represent a finite number of digits.
- Operations between these finite representations of numbers introduce round-off error, look at how numbers are stored.

Integers.

Look at an 8-bit word

There are 8 pieces of base-2 bits
2 values of 0 or 1



0 \Rightarrow positive
1 \Rightarrow negative

Other 7 bits give the #
in base-2

ex.) 1 | 1 0 0 1 0 1 |

$$\Rightarrow -(2^6 + 2^3 + 2^1 + 2^0) = -75$$

Thus, the smallest # is:

$$0000000_{\text{base } 2} = 0_{\text{base } 10}$$

$$1111111_{\text{base } 2} = 127_{\text{base } 10}$$

with the first bit: -127 to 127

Since $+0 = -0$

\Rightarrow Use -0 to represent -128

\Rightarrow 8-bit can represent $2^8 = 256$ pieces of data $(-128 \rightarrow 127)$

- Integer Arithmetic is exact so long as no remainder!

$$\frac{8}{2} = 4 \quad \frac{7}{2} = 3$$

- You can get under or over flow,

Underflow: $-74 \times 2 = -148 \leftarrow$ outside of $-128 \rightarrow 127$

Overflow: $123 + 45 = 168 \leftarrow$ outside of $-128 \rightarrow 127$

In modern Computers: 32-bit or 64-bit


32-bit: $2^{31} = 2,147,483,648$

64-bit: $2^{63} = 9,223,372,036,854,775,808$


Side note! This is the reason that old computers could not use more than 4 Gb of memory, 2 Gb for OS + 2 Gb other.

Floating Point Representation

Real numbers (called floating point) are stored via the IEEE 754 specification.



Sign is 1 or 0 for negative or positive
Signed exponent is the maximum value of the base
Mantissa contains the significant digits.



$$N = \pm d_1 d_2 d_3 \dots d_p B^e = m B^e$$

$m =$ mantissa

$B =$ base of the # system

$e =$ "signed" exponent

ex. 1) 8-bit base-10 representation

$$1 \mid 095 \mid 1467 \quad B=10$$

$$\text{mantissa: } m = -(1 \times 10^{-1} + 4 \times 10^{-2} + 6 \times 10^{-3} + 7 \times 10^{-4}) \\ = -0.1467$$

$$\text{Signed exp: } e = +(9 \times 10^1 + 5 \times 10^0) = 95$$

$$10951467_{\text{base } 10} = m B^e = -0.1467 \times 10^{95}$$

Normalization

$$1 \text{ in}^2 = \frac{1}{144} \text{ ft}^2 = 0.\overbrace{006944}^{\text{need to store}} \text{ ft}^2$$

This is less accurate than

$$1 \text{ in}^2 = \underbrace{0.694444}_{\text{need to store}} \times 10^{-2} \text{ ft}^2$$

- Remove leading zeros by lowering the exponent.

In general you want the mantissa to be bounded by:

$$\frac{1}{B} \leq m < 1 \rightarrow \text{base-10} \quad \frac{1}{10} \leq m < 1$$

$$\text{base-2} \quad \frac{1}{2} \leq m < 1$$

If $m < 1/2$, multiply by 2, increase exponent

Single Precision

A real number is stored in
4 bytes (for 32 bit)

bit (binary digit): 0 or 1

1 byte = 8 bits

Word = 4 bytes (32-bit), 2 bytes (64-bit)

32 bit $\left\{ \begin{array}{l} 23 \text{ for digits} \\ 8 \text{ for signed exponent} \\ 1 \text{ for sign} \end{array} \right.$

Max exponent: $\pm 2^7 = \pm 128$

64 bit (Double Precision)

64 bits $\left\{ \begin{array}{l} 52 \text{ for the digits} \\ 11 \text{ for signed exp} \\ 1 \text{ for sign} \end{array} \right.$

Signed exponent: $\pm 2^{10} = \pm 1024$

Single vs. Double

Smallest

Largest

Single $\pm 1.0 \dots \times 2^{-126} \approx \pm 1.1 \times 10^{-39}$

$\pm 1.11 \times 2^{127} \approx \pm 3.4 \times 10^{38}$

Double $\pm 1.0 \dots \times 2^{-1022} \approx 2.22 \times 10^{-308}$

$\pm 1.11 \times 2^{1023} \approx 1.79 \times 10^{308}$

Note: You can't access 10^{-128} or 10^{-1024} because those are reserved for infinity / NaN

Matlab uses Double precision.

Round off error,

Round off error is when you do not have enough bits to store the answer,

In double precision, the precision is 2^{-52} (as determined by the mantissa)

This is roughly $2.22 \times 10^{-16} \leftarrow$ called machine precision & given by eps in matlab

Simplified example of Round off error,

Addition Problem: $0.99 + 0.0044 + 0.0042 = 0.9986$ (exact)

3-digit arithmetic:

$$(0.99 + 0.0044) + 0.0042 = 0.994 + 0.0042 = 0.998$$

$$0.99 + (0.0044 + 0.0042) = 0.99 + 0.0086 = 0.999$$

There are ways to combat this.

Cancellation error:

Look at $x^2 - bx + 1 = 0$, b is large,
 r close to b

$$r = \sqrt{b^2 - 4}$$

$$x_1 = \frac{b+r}{2}$$

$$x_2 = \frac{b-r}{2}$$

If b & r are
close \rightarrow potential
for errors.

Reformulate

$$x_2 = \frac{b-r}{2} \cdot \frac{b+r}{b+r} = \frac{b^2 - r^2}{2(b+r)} = \frac{4}{2(b+r)} = \frac{2}{b+r}$$

ex.) let $b = 97$, $r = 96.9794$

$$x^2 - 97x + 1 = 0$$

exact: 0.01031

Standard: 0.01050 (3-digit math)

reformulate: 0.01031